

Instructions for MA2

This document contains the material for the second module (MA2) in the course Computer Programming II. This module is an introduction to data structures using Python. We introduce some aspects about classes and other topics like *operator overloading*, *iterators* and *generators*.

The exercises for the module can be found in the yellow frames below. There are **14 mandatory exercises** in this module. Each group of exercises follows a related theory and a set of examples.
For this module, feel free to write classes from scratch.

The following files are associated with the module:

1. `linked_list.py` in which you shall write the solutions to linked list exercises,
2. `bst.py` in which you shall write the solutions to binary tree exercises, and
3. a number of test files with the name `test_LinkedList_name` or `test_BST_name` where `name` is the name of the class that is tested, `LinkedList` and `BST`, respectively.
4. There is a file `tests.sh` to run all tests from the terminal on some operating systems.

Follow the instruction below when preparing your solutions for presenting:

1. The tasks must be solved in the order they come and with the tools (classes, functions, methods, ...) which are introduced before the exercise. The point of this exercise is not only to achieve a specific result, but to do it using specific techniques and approaches.
2. Unless otherwise stated, you may not use other packages or pre-written code, than what is already included in the downloaded files.
3. The functions must be named, take arguments, and return exactly what is stated in the task.
4. Your code may not use any global variables.
5. Answer text tasks by including comments or printed text strings at the end of the code document. The answers can be kept succinct.
6. Present your solutions to a teacher or assistant in the class or a Zoom room. Be prepared to answer questions about the complexity of the methods.
7. When your orally presented solutions are approved:
 - (a) Fill in the name, email, assistants / teachers who reviewed the assignment and the date of the review the documents.
 - (b) Make a zip file containing `linked_list.py` and `bst.py` as well as test files and upload into STUDIUM.

Important note:

You may collaborate with other students, search the Internet, and get a general overview using AI tools, but you must write and be able to explain the code you actually present. You may not copy code, either from other students, automated tools, or from the Internet, except from those places explicitly pointed out in this module.

Changing variable names and similar superficial modifications does not count as writing your own code. Also, aiding somebody to cheat (for example, if somebody hands in your code as their own) is not permitted and will be reported.

Since these assignments are included as part of the course examination, we are obliged to report blatant violations of these rules to the disciplinary board.

All course material that we provide is copyrighted and may not be spread.

1 Data structures

A *data structure* is a way of organizing a collection of values. Examples of such collections include a car register, a family tree, a table with distances between different places etc. We call the individual objects (i.e. cars, people, places, ...) *records* or, sometimes, **elements**.

We will also discuss different operations on the structures and analyze their efficiency. Examples of such operations:

- Insert new records.
- Remove records.
- Search for records with specific contents.
- Visit all records.

Some common data structures:


- Arrays.
- Linear linked lists.
- Trees.
- General linked structures (graphs).
- Hash tables.

Arrays

This is the most fundamental data structure since it is closely related to the actual memory organization in computers.


Typical features of arrays:

- all elements are of the same data type (e.g. integer) which means that they take up the same amount of space and
- all elements are stored in a contiguous memory area.



These properties mean that the address of any element can be easily calculated if you know the address of the first element (in place 0) and the element size.

If the elements are of different sizes, we can store pointers to the elements (which gives us an array of pointers)




The main advantage of arrays is the fast index operation ($\Theta(1)$).

A disadvantage is that it is more difficult to add a new element at a given index because all subsequent elements must be moved ($\Theta(n)$).

Linear linked lists


Instead of storing the records in a contiguous memory, we can let them point to each other i.e. a record can contain *the address of* the next record.



The advantage of this is that it is possible to link in new elements anywhere without moving already stored elements. One drawback is that the index operation requires $\Theta(n)$ operations because we have to start from the beginning and follow pointers until we come to the element with the sought index.


Trees

In linked lists, each record has (at most) one successor. If you allow a record to have multiple successors, we get a *tree structure*:



Graphs

If a record can have several relations both backwards and forwards and also circularly, we call the structure a graph:



Hash table

A hash table is a way of implementing a mapping from a set of keys to a set of values such that each element in the key set relates to one specific element in the value set.

Data structures in Python

The basic data structure in Python is the *list*. A list is implemented as an array of references to the records, thus allowing fast indexing ($\Theta(1)$). Since lists can be expanded, the array must be *dynamic* i.e. it can grow when needed.

The *tuple* data type is similar to lists, but since it is unchangeable (references cannot be updated or added), they are sometimes more efficient.

The *dict* data type represents pairs of keys and values. It is implemented using hashing techniques.

Many packages and modules define other data structures. For example, the NumPy package defines *multidimensional arrays* and the pandas package defines *dataframes*.

2 Linked lists


Python's built-in, array-based, lists are extremely powerful and often the best option for general lists. There are still reasons to study how to make linked lists. It's an easy way to illustrate techniques that handle all sorts of dynamic structures — not just linear lists.

In addition, it is a good exercise in recursive thinking which can be applied to more complicated structures such as trees and graphs.

We will exemplify the technique with lists of integers as data in the nodes. We keep the lists sorted so that the elements come in ascending order.

Using integers as data imposes no limitation. The only thing that is important is that the elements are comparable with less than/greater than.

A linked list with integers can be illustrated as follows:




We represent the individual elements in the list as objects from the class `Node`:

```
1 class Node:
2     def __init__(self, data, succ):
3         self.data = data
4         self.succ = succ
```

(In this situation we don't care about "protecting" i.e. to prefix supposedly internal variables name the names with `_` for the variables in the class.)

To keep track of the list, you need a reference to the first element.

Here is an illustration of the effect of three lines of code:




Now we can iterate over the list and, for example, sum the stored numbers:

```
1 summa = 0
2 f = first
3 while f:
4     summa += f.data
5     f = f.succ
6 print('Sum: ', summa)
```

The code will produce the line `Sum: 25`

Now suppose we want to write code that adds a new value to a list keeping it sorted.




There are four different cases:

1. In an empty list.
2. As the first element. For example inserting **1**
3. As the last element. For example inserting **53**
4. In the inner. For example inserting **8**

We have to find the pointer to be changed, i.e. either the pointer to the first node or the

pointer in the node that will be *immediately in front of* the new node:



Assume that the value we want to insert is in the variable `x`. The code can then be written:

```

1 if first == None or x <= first.data:    # case 1 and 2
2     first = Node(x, first)
3 else:
4     f = first
5     while f.succ != None and x > f.succ.data:
6         f = f.succ
7     f.succ = Node(x, f.succ)           # case 3 and 4
8

```

In case the value already exists, we store the new value in front of the old value. If we want to store the new value after the old value, the code becomes a little more complicated. (If only integers are stored, it is unimportant how equal numbers are stored *but* if it is an object with several properties then the insertion order can be important.)

It is a special case if the new value is to be placed first in the list, due to list being empty or due to the new value being less than the first element. In those cases, the variable `first` must be changed.

If the new value is not to be the new first element, we use a `while`-loop to find the node that should come *immediately in front of* the new node.

Remember that the order is important — you must first know that there is a node present before you can look at its contents!

A class for the list

It is convenient to put the linked list machinery in a separate class. Sketch:

```
1 class LinkedList:
2     class Node:
3         def __init__(self, data, succ):
4             self.data = data
5             self.succ = succ
6
7     def __init__(self):
8         self.first = None
9
10    def print(self):    # Unnecessary but serves as a simple example
11        pass
12
13    def insert(self, data):
14        pass
15
```

The class `Node` is the same as before, but it is now placed inside the class — it has no use outside `LinkedList`.

The class `LinkedList` has only one instance variable: the first node (or, more precisely, a reference to the first node). Initially, it is set to `None` i.e. the list is empty.

Insertion in the list

We place the code for insertion, which was previously described, in a method:

```
1 def insert(self, x):
2     if self.first is None or x <= self.first.data:
3         self.first = LinkedList.Node(x, self.first)
4     else:
5         f = self.first
6         while f.succ and x > f.succ.data:
7             f = f.succ
8         f.succ = LinkedList.Node(x, f.succ)
```

Since the class `Node` is an attribute in the class, we must refer to it with `LinkedList.Node`. Here we write `f.succ`, rather than `f.succ != None` or `f.succ is not None`, as the first condition in the loop on line 6. In a condition expression, `None` is interpreted as `False` so this is just a shorter way to write the same thing. This is only safe if any non-`None` value will be interpreted as `True`. Since we only store `Node` instances, that holds.

A simple print method

As an example of a simpler method than the `insert` method, we show a method for printing lists. It is actually unnecessary because printing could be accomplished by implementing the magical method `__str__`, but it serves as a simple example.

First attempt:

```
1 def print(self):
2     f = self.first
3     while f:
4         print(f.data)
5         f = f.succ
```

This gives one value per line, which is probably not the way you want to see the list.

With a little more effort, we can get all values on the same row, separated by comma and the whole list enclosed by parentheses:

```
1 def print(self):
2     print('(', end='')
3     f = self.first
4     while f:
5         print(f.data, end='')
6         f = f.succ
7         if f:
8             print(', ', end='')
9     print(')')
```

The code that we have discussed so far is in the downloaded `linked_list.py`.

Exercise 1: Method `length`

Write a method `length(self)` which, using a `while`-loop, calculates and returns the number of nodes in the list. Original list must not be destroyed.

Exercise 2: Method `remove_last`

Write a method which removes the last node from the list. The method should return the value in the removed node.

If the list is empty, a `ValueError` should be raised. (`ValueError` is a standard exception class in Python.)

Exercise 3: Method `remove`

Write a method `remove(self, x)` which deletes the first node containing `x` as data. If the method finds a node with this content, it should return `True` else `False`.

Use iteration here. We will discuss how to do this recursively soon. Make sure that the method works even if the element to be removed is currently the first or the last element of the list.

3 Recursive list methods

Linked structures are well-suited for recursive methods. This is because a list can be defined recursively in this way:

A *list* is
either
empty (consists of 0 elements)
or
consists of *one element followed by a list*.

We will now show how to implement the method `length` recursively, rather than the iterative approach used in the exercise above. This can be done in two ways.

Approach 1: With a helper method in the `Node` class

We write a method in the *node class* which tells you the length of the sublist starting at that node:

```
1 class Node:
2     ...
3     def length(self):
4         if self.succ is None:
5             return 1
6         else:
7             return 1 + self.succ.length()
8     ...
```

Thus, if this node has no successor, the length is 1. Otherwise it is 1 *plus* the length of the list beginning with the successor.

Then, in the class `LinkedList` we write the required method:

```
1 def length(self):
2     if self.first is None:
3         return 0
4     else:
5         return self.first.length()
```

We have to start by checking if there are any nodes in the list at all.

Approach 2: Using a helper method in the `LinkedList` class

In the second approach, we do nothing in the `Node` class. Instead, we write an internal method

`_length(self, f)` in the `LinkedList` class that returns the number of nodes in the sublist starting at `f`. Then we use this method in the externally callable method:

```
1 def length(self):
2     return self._length(self.first)
3
4 def _length(self, f):
5     if f is None:
6         return 0
7     else:
8         return 1 + self._length(f.succ)
```

The helper method can also be placed as a local function within `length`.

```
1 def length(self):
2
3     def _length(f):
4         if f is None:
5             return 0
6         else:
7             return 1 + _length(f.succ)
8
9     return _length(self.first)
```

Since it is now not a method in the class, it has no `self` as a parameter or `self.` before the call.

The choice between the two approaches is a matter of taste. Since the second approach often gives fewer special cases, the code becomes shorter, especially when it comes to trees. However, these approaches may be more difficult to understand.

Recursive `insert` using approach 1

Here we will write an `insert` in the `Node` class. Within a method in the `Node` class, it is not possible to change the pointer *to* that node but only the pointer *in* that node, to its successor. Just as in the iterative variant, we must therefore stop at the node which is to be immediately before the new node. This means that we should only proceed in the recursion if the current successor node is also supposed to be located before the new node. Thus, we get the following code in the `Node` object:

```
1 def insert(self, x):
2     if self.succ is None or x <= self.succ.data:
3         self.succ = LinkedList.Node(x, self.succ)
4     else:
5         self.succ.insert(x)
```

The new node is created on line 3.

The method in the main class must first check that the new node is not in the first place, i.e. the code must ensure that there is at least one element in the list, and that this element should come before the new node.

This amounts to the following code added within `LinkedList`:

```
1 def insert(self, x):
2     if self.first is None or x <= self.first.data:
3         self.first = LinkedList.Node(x, self.first)
4     else:
5         self.first.insert(x)
```

It feels a bit ugly to have to repeat basically the same code. ...

Recursive `insert` using approach 2

In approach 2, we place a helper method in the main class. The helper receives a reference to the first node and returns the intended first node in the modified list.

```
1 def insert(self, x):
2     self.first = self._insert(x, self.first)
3
4 def _insert(self, x, f):
5     if f is None or x <= f.data:
6         return LinkedList.Node(x, f)    # Return the new node
7     else:
8         f.succ = self._insert(x, f.succ)
9         return f                        # Return the same node we got in
```

Note how the helper method `_insert` is called on line 2! The new node can become the first in the list.

The helper works as follows:

It gets a reference to the first node in a sublist and returns the first node in the modified sublist. If the new node has become the first in the sublist, that is the one which is returned, but if it ended up further away then the "old" first node is returned (line 9).

The nice thing about this method is that we do not have to handle a new first node or an empty list as special cases.

The next exercise should be solved with recursion rather than iteration. Go on working with the code you downloaded earlier.

Exercise 4: The method `to_list`.

Write a recursive method `to_list (self)` which returns a standard Python list with the values from the linked list in the same order.

3.1 Iterators and generators for the class `LinkedList`

To process all elements in many of the common structures in Python (lists, tuples, strings, files ...) we can use a `for` statement. It is said that these structures are *iterable*. For a structure to be iterable, there must be a method `__iter__` which returns an *iterator*.

An iterator is an object which keeps track of a *current element*. The class must include one specific method, `__next__`. Our class is supposed to both be iterable and an iterator.

An iterator for objects of class `LinkedList` could look like this:

```
1 class LinkedListIterator:
2     def __init__(self, lst):
3         self.current = lst.first
4
5     def __iter__(self):
6         return self
7
8     def __next__(self):
9         if self.current:
10            res = self.current.data
11            self.current = self.current.succ
12            return res
13        else:
14            raise StopIteration
```

The constructor (the method `__init__`) receives (a reference to) a `LinkedList` object and saves a reference to its first node in the instance variable `current`. The method `__next__` returns the value in the current element and advances `current` to the next item in the list.

If the last element has been passed, the exception `StopIteration` is raised.

The method `__iter__` just returns itself.

If this class exists, we can for example write:

```
1 it = LinkedListIterator(lst)
2 for x in it:
3     print(x)
```

So the `for`-statement in Python uses exactly these methods. `__iter__` is called first, and then `__next__` is called repeatedly on the returned object.

You can also iterate over the list with a `while` statement which will rely on `__next__` in `LinkedListIterator`:

```
1 it = LinkedListIterator(lst)
2 while True:
3     try:
4         print(next(it))
5     except StopIteration:
6         break
```

This gives an idea of how the `for` statement is implemented.

Note that `next(it)` is a nicer way to write `it.__next__()`. If you want to try this iterator you can download [linked_list_iterator.py](#) (optional).

Instead of making a separate iterator class, you can provide the class `LinkedList` with the methods `__iter__` and `__next__`.

The method `__iter__` is then given the task of creating the `current` variable:

```
1 def __iter__(self):
2     self.current = self.first
3     return self
```

Note that this method thus adds an instance variable to the class!

The method `__next__` is completely unchanged.

Now we can iterate over linked lists in the same way as we iterate over ordinary lists:

```
1 for x in lst:
2     print(x)
```

There might be drawbacks with this approach. What if we wanted to iterate over all pairs:

```
1 for x in lst:
2     for y in lst:
3         print(x, y)
```

This can be achieved even easier by making `__iter__` as a *generator*:

```
1 def __iter__(self):
2     current = self.first
3     while current:
4         yield current.data
5         current = current.succ
```

The Python statement `yield` will act as a `return` statement *but* next time the generator is called, it will continue where it left off. We do not need to write any `__next__` method or throw any `StopIteration` — these are created automatically.

This generator is included in the code you downloaded from the beginning. Will the iteration over pairs above work with this one? Why?

More about iterators and generators can be found on YouTube, for example in [iterators](#) and [generators](#) by Corey Schafer.

3.2 Operator overloading

For standard lists we can use the operator `in` and, for example, write expressions like `if x in list:`. We can also introduce this in our own list class by defining the method `__contains__`:

```
1 def __contains__(self, x):
2     for d in self:
3         if d == x:
4             return True
5         elif x < d:
6             return False # No point in searching more
7     return False
```

This code is included in the class you downloaded in the beginning.

Note how we can use our own class generator in the `for` statement.

In a similar way, we can define other operators such as `<`, `<=`, `==`, ... by defining the methods `__lt__`, `__le__`, `__eq__`, ...

More about magic methods and operator overloading can be found, for example, in Corey

Exercise 5: The method `__str__`.

Use the iterator or generator to write the method

`__str__(self)` which returns the list as a string of comma-separated values with parentheses surrounding the list.

Examples of resulting strings: `()`, `(1)`, `(1, 2)`. Note that the result must have *exactly* that appearance, i.e. with commas followed by spaces *between* every value.

Exercise 6: The method `copy`.

The `copy` method in the downloaded code looks like the code to the right:
It creates and returns a copy of the `LinkedList` object.

```
1 def copy(self):  
2     result = LinkedList()  
3     for x in self:  
4         result.insert(x)  
5     return result
```

What is the complexity of this method?

Rewrite the code to improve its complexity! Note that the copy may not share any list nodes with the original!

What is the complexity of your method?

Exercise 7: A `Person` class in `LinkedList`

Suppose you have the class:

```
1 class Person:  
2     def __init__(self, name, pnr):  
3         self.name = name  
4         self.pnr = pnr
```

1. Create an empty list `plist` in `main` and then a person object `p`.
2. Insert the person object into the list using `insert` method of the `LinkedList`. Print `plist`. This should work well.
3. Create another person object `q` and insert it with `insert`.
The program will be aborted in the `insert` method on line
`if f is None or x <= f.data:`
The diagnosis is
`TypeError: '<=' not supported between instances of 'Person' and 'Person'`

The problem is that Python doesn't know how to compare two person objects in terms of size — `insert` is supposed to make sure that the list is sorted in order of size.

To be able to use relational operators such as `<`, `<=`, `==` ... between two objects the methods `__lt__`, `__le__`, `__eq__`, ... must exist.

Write the methods in the `Person` class that are needed to be able to insert such objects into the `LinkedList`.

Note that you *not* should change anything in the `LinkedList` class but define comparison operators in the class `Person`.


Choose whether `Person` objects should be sorted by name or social security number (`pnr`).

4 Tree structures

Definitions and terminology

In a list, each node has exactly one successor (except the last) and one predecessor (except the first).

If we allow a node to have multiple successors, we get a *tree structure*:



The terminology is taken from both the forest and the family:


A *parent* is a node that has one or more successors (*children*).

A *child* is a node that has a predecessor (*parent*).

Nodes with the same parent are called *siblings*. The siblings come in order of age with the oldest first (furthest to the left).

A *leaf* is a node with no children.

All nodes except the *root* have a parent.




The tree *height* is the largest number of nodes on the way between the root and a leaf. The height of the above tree is thus 4.

Sometimes the height is defined as the number of ‘arcs’ between the root and a leaf. With this definition, the height in the tree above is 3. It does not matter much which definition you use, as long as you are consistent.

A *binary tree* is a set of nodes that is either empty or consists of three subsets:

1. one with just one node – the *root*,
2. one subset called the *left subtree* and
3. one subset called the *right subtree*.

Furthermore, the two sub-trees are in turn also binary trees.



Note the difference between a binary tree and an ordered tree with a maximum of two children per node:

1. An ordered tree has at least one node while a binary tree can be empty.

2. In a binary tree, there is a difference between having an empty right subtree and an empty left subtree — no such difference exists for ordered trees.

Tree operations

There are a number of operations you may want to perform on trees depending on the application:

- *Traverse* i.e. visit all the nodes (in some order).
- *Search* for a node with a certain content or in a certain position.
- *Insert* and *delete* nodes.
- *Merge* trees.
- Compute various metrics (height, number of nodes, path lengths, level of balancing, ...).

Here we only discuss algorithms for traversals.


Other operations are discussed in the coding examples and the assignments.

Tree traversals

Traversing a structure means visiting all the elements in some particular order. For trees, there are mainly four possible schemes: *preorder*, *inorder*, *postorder* and *level order (breadth first)*.

Preorder:

1. Visit the root
2. Visit the children in *preorder* (with the oldest first)



On the example tree, the order will be: e b a d c n g f k p y

Postorder:

1. Visit the children in *postorder* (with the oldest first)
2. Visit the root

On the example tree, the order will be: a c d b f k g y p n e

Inorder (for binary trees):

1. Visit the left subtree in *inorder*
2. Visit the root
3. Visit the right subtree in *inorder*

On the example tree, the order will be: a b c d e f g k n p y

Level order:

1. Visit the root (level 0)
2. Visit all nodes on level 1
3. Visit all nodes on level 2
4. ...


On the example tree, the order will be: e b n a d g p c f k y

Binary search trees


Binary *search trees* are an important application of binary trees. In such a tree, each node has a *key* which uniquely identifies the node. A key can be, for example, a social security number, a vehicle registration number or a word. Typically, there are no two nodes with the same key.

The keys must have an ordering defined, i.e. one should be able to compare the keys in size.

The key in a node must be larger than all keys in the node's left subtree and less than all keys in the node's right subtree:




Example: A binary search tree with integers as keys:




To search for a node with a given key, we start from the root. The order relation can then guide the search down to the requested node.

For example, to find **13**, go to the right at **12**, to the left at **23** and **17** since **13** is greater than **12** but less than both **23** and **17**:



The same is done when inserting new keys: start from the root and let the nodes guide the direction. There is always exactly one place that a new key should go into. The figure shows insertion of the key **14**:



The work of both searching and inserting keys in a binary search tree has an upper bound defined by *the height* of the tree.

Representation and algorithms for binary search trees

Here we will build a structure for storing nodes in binary search trees. In the code, we name the stored records **keys** and use integers as an example, but all types of data that have the ordering relations defined work well. We do not allow the same key to occur multiple times.

We want to be able to create trees, add new records, search for certain records and print the contents of the tree.

```

1 class BST:
2
3     class Node:
4         def __init__(self, key):
5             self.key = key
6             self.left = None
7             self.right = None
8
9     def __init__(self):
10         """Create an empty tree"""
11         self.root = None
12
13     def insert(self, key):
14         """Insert a new key in the tree"""
15         pass
16
17     def contains(self, k):
18         """Check if k is in the tree"""
19         pass
20
21     def print(self):
22         """Print the contents of the tree"""
23         pass
24

```

The method `print` is included from the beginning to be able to easily print trees but it will later be replaced by the more general `__str__` method.

The tree thus consists of node objects where each node contains a key as well as references to the left and right subtrees.

The tree object only has the instance variable `root` which is initialized to `None` i.e. the tree is empty from the beginning.

We start by looking at the method `contains`:

Note that the search is controlled by the relationship between the searched value and the contents of the node.

The iteration process continues as long as `n` is not `None` and the key is not in `n`. If `n` is not `None` when the loop ends, the value has been found.

```

1     def contains(self, k):
2         n = self.root
3         while n and n.key != k:
4             if k < n.key:
5                 n = n.left
6             else:
7                 n = n.right
8         return n is not None

```

We have written this method iteratively. In many cases, however, it is much easier to write tree methods recursively. The `print` method below is an example of this.

The method prints the tree on one line with the nodes sorted by size.

```

1     def print(self):
2         self._print(self.root)
3
4     def _print(self, r):
5         if r:
6             self._print(r.left)
7             print(r.key, end=' ')
8             self._print(r.right)
9

```

This pattern with an auxiliary method (which gets a reference to a node to work with) and one main method (which only starts the helper by passing the root node) will be used in many tree methods.

Unlike the `print` method, the `insert` method is easy to write iteratively, but we do it recursively anyway.

```
1  def insert(self, key):
2      self.root = self._insert(self.root, key)
3
4  def _insert(self, r, key):
5      if r is None:
6          return self.Node(key)
7      elif key < r.key:
8          r.left = self._insert(r.left, key)
9      elif key > r.key:
10         r.right = self._insert(r.right, key)
11     else:
12         pass # Already there
13     return r
```

We also showed this technique as “approach 2” in the section on linked lists, i.e. the helper method *returns* a reference to the root node in the modified subtree. The nice thing about this technique is that we do not have to explicitly handle separate cases for a completely empty tree, the new node being a left child and the new node being a right child.

We show another method which calculates the number of nodes. It uses the same recursive technique with a helper:

```
1  def size(self):
2      return self._size(self.root)
3
4  def _size(self, r):
5      if r is None:
6          return 0
7      else:
8          return 1 + self._size(r.left) + self._size(r.right)
```

A generator

In the `LinkedList` class we presented a generator to be able to iterate over the nodes. The keyword there was `yield`. We will now look at a generator for our BST class which goes over the nodes in *inorder*, meaning order of magnitude.

We start by writing a generator for the `Node` class:

```

1  class Node:
2      ... # Same as before
3      def __iter__(self):
4          if self.left:          # If there is a node to the left
5              for n in self.left: # use its generator
6                  yield n.key    # to return node references
7          yield self             # Then return ourself
8          if self.right:         # Then the nodes to the right
9              for n in self.right:
10                 yield n
11

```

It is recursive in the sense that a generator in a node uses a generator in other nodes.

Note that we need to check that there is a node before we try using its generator.

This can be written a little more compactly using the construction `yield from`:

```

1  class Node:
2      ...
3      def __iter__(self):
4          if self.left:
5              yield from self.left
6          yield self.key
7          if self.right:
8              yield from self.right
9

```

Now we can easily make a generator for the tree:

```

1  def __iter__(self):
2      if self.root:
3          yield from self.root

```

Deletion of nodes in binary search trees

Removing a key from a search tree is a bit more contrived than the other operations we've covered. If the key is in a node with no or just one child, there is no problem. It is basically the same thing as removing nodes from a linked list. The problem is removing keys that are in nodes with two children because we then do not know how to keep track of the node's children.

The usual technique is to keep the physical node but replace the content (key) with another key that is to be kept in the tree. Since we have to maintain the search tree property (small on the left, large on the right), there are only two possible keys to move up to the node where the key to be removed is: the *largest* among the small ones or the *smallest* among the large. Both of these have the property of having at most one child — otherwise they are not the largest or the smallest. Which of them you choose does not matter.

Here, we give a pseudocode for the operation:

```

1 def remove(self, key):
2     self.root = self._remove(self.root, key)
3
4 def _remove(self, r, k):
5     if r is None:
6         return None
7     elif k < r.key:
8         r.left = # left subtree with k removed
9     elif k > r.key:
10        r.right = # right subtree with k removed
11    else: # This is the key to be removed
12        if r.left is None: # Easy case
13            return r.right
14        elif r.right is None: # Also easy case
15            return r.left
16        else: # This is the tricky case.
17            # Find the smallest key in the right subtree
18            # Put that key in this node
19            # Remove that key from the right subtree
20    return r # Remember this! It applies to some of the cases above

```

Note that there is recursion on lines 8, 10 and 19.

Exercise 8: A method `contains`

Write the `contains` method with recursion instead of iteration.

Note that `contains` is checked in the `main`.

There is no test file for this function, but you may first write suitable tests that are satisfied by the existing function `contains`. Then, you can apply the same tests on your own function.

Exercise 9: Compute the height

Write a method `height` which returns the height of the tree. The empty tree has height 0, the tree with 1 node has height 1, the tree with 2 nodes has height 2 while the tree with 3 nodes have a height of 2 or 3 depending on the order of insertion.

Exercise 10: A method `__str__`

The generator code is included in the downloaded file. Use it to write a `__str__` method! The elements in the string should be *separated* by commas followed by a space and the string should be surrounded by '`<`' and '`>`'.

Examples of results: '`<>`', '`<1>`' and '`<1, 2>`'.

Exercise 11: A method `to_list`

Write the method `to_list` which creates and returns a standard Python list with the values from the tree. What is the complexity?

Exercise 12: A method `to_LinkedList`

Write the method `to_LinkedList` that creates and returns a `LinkedList` with the values from the tree. Complexity of your algorithm should be linear.

Exercise 13: Remove a key

Complete the method `remove` according to the given pseudocode.

Exercise 14: Complexity of search

In a binary search tree with n nodes and height h , what is the complexity in the following scenarios:

- Worst case of successful search
- Worst case of failed search

Later voluntary exercises will cover more theory regarding the complexity of binary tree operations.